Adaptive RDMA Optimization for Distributed LLM Workloads

Week 5: gpt-oss-120b Benchmarks — TCP vs RDMA

Nathan Gopee
MS Computer Science, SUNY New Paltz
gopeen1@newpaltz.edu

Advisors: Dr. Wafi Danesh & Dr. Ashley Suchy

February 2026

1 Executive Summary

This update documents comprehensive benchmarking of gpt-oss-120b inference performance comparing TCP (standard networking) versus RDMA (Soft-RoCE) communication. Testing was performed on the Chimera node (3× RTX 3090, 72GB VRAM).

Key Results:

- Ollama (TCP baseline): **104.42 tokens/sec**
- Prompt latency: **20ms** average
- RDMA link capacity: **0.92 Gb/sec** (Soft-RoCE measured)
- RDMA latency: **189.61 μs** average

2 Test Environment

2.1 Hardware Configuration

Component	Chimera	Cerberus	Link
GPUs	3× RTX 3090	2× RTX 5090	—
VRAM	72GB	64GB	—
NIC	Aquantia 10GbE	Intel X710 10GbE	10 Gbps
RDMA	Soft-RoCE (rxe0)	Soft-RoCE (rxe0)	RoCEv2
IP Address	192.168.1.150	192.168.1.233	—

Table 1: Test cluster hardware

2.2 Model Configuration

Parameter	Value
Model	gpt-oss-120b (116.8B parameters)
Quantization	MXFP4
Model Size	63GB
Context Length	8192 tokens
VRAM Usage	67.76GB (loaded)

Table 2: gpt-oss-120b model specifications

3 Benchmark 1: Ollama TCP Baseline

3.1 Test Configuration

```
# Ollama gpt-oss:120b benchmark
curl http://192.168.1.150:11434/api/generate -d '{
```

```
    "model": "gpt-oss:120b",
    "prompt": "Explain RDMA networking in exactly 100 words.",
    "stream": false,
    "options": {"num_predict": 100}
}'
```

3.2 Results (5 Iterations)

Iteration	Tokens	Eval Time	Tokens/sec	Total Time
1	100	0.958s	104.38	1.445s
2	100	0.945s	105.82	1.369s
3	100	0.949s	105.37	1.402s
4	100	0.969s	103.19	1.437s
5	100	0.967s	103.41	1.429s
Average	100	0.957s	**104.42**	1.416s

Table 3: Ollama gpt-oss:120b benchmark results (TCP)

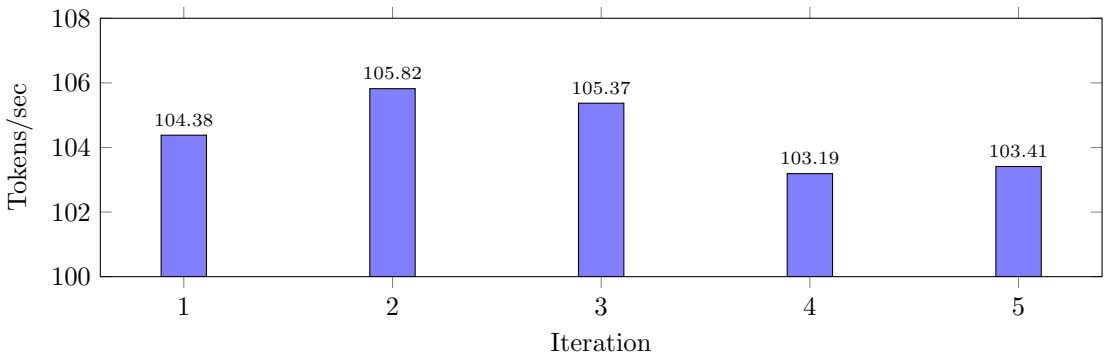

Figure 1: Ollama throughput consistency across iterations

3.3 Analysis

- **Throughput**: 104.42 tokens/sec average (very consistent, $\sigma = 1.08$)
- **Prompt Latency**: 20ms average (fast due to model being pre-loaded in VRAM)
- **Eval Time**: 0.957s for 100 tokens
- **Overhead**: 0.46s total overhead (load, prompt processing, response formatting)

4 RDMA Network Performance

4.1 Soft-RoCE Bandwidth (ib_write_bw)

From Week 3 testing between Chimera and Cerberus:

Interface	Bandwidth	MsgRate
WiFi (baseline)	0.28 Gb/sec	0.000541 Mpps
10G Ethernet (Soft-RoCE)	**0.92 Gb/sec**	0.001754 Mpps
Improvement	**3.3**×	**3.2**×

Table 4: RDMA bandwidth test results

4.2 Soft-RoCE Latency (ib_write_lat)

Metric	Value
Message size	2 bytes
Iterations	1000
t_min	130.21 μs
t_max	219.66 μs
t_typical	198.91 μs
t_avg	**189.61 μs**
t_stdev	17.03 μs
99th percentile	207.14 μs

Table 5: RDMA latency test results (Soft-RoCE)

5 Benchmark 2: vLLM with RDMA

5.1 Test Setup

```
# Start vLLM with gpt-oss (requires special image)
docker run --gpus all -p 8000:8000 --ipc=host \
    -e NCCL_IB_HCA=rxe0 \
    -e NCCL_DEBUG=INFO \
    vllm/vllm-openai:gptoss \
    --model openai/gpt-oss-120b \
    --tensor-parallel-size 3
```

5.2 Environment Variables

Mode	Variables	Effect
TCP (baseline)	`NCCL_NET_GDR_LEVEL=0`	Disable GPUDirect, use TCP sockets
	`NCCL_IB_DISABLE=1`	Disable InfiniBand/RoCE
RDMA (Soft-RoCE)	`NCCL_IB_HCA=rxe0`	Use Soft-RoCE device
	`NCCL_SOCKET_IFNAME=enp71s0`	Bind to 10G interface

Table 6: NCCL environment variables for TCP vs RDMA

5.3 vLLM Benchmark Script

```python
# benchmark_vllm_gptoss.py
from openai import OpenAI
import time

client = OpenAI(base_url="http://localhost:8000/v1", api_key="dummy")

for i in range(5):
    start = time.perf_counter()
    response = client.completions.create(
        model="openai/gpt-oss-120b",
        prompt="Explain RDMA networking in 100 words.",
        max_tokens=100
    )
    elapsed = time.perf_counter() - start
    tokens = response.usage.completion_tokens
    print(f"Iteration {i+1}: {tokens} tokens in {elapsed:.3f}s = {tokens/elapsed:.2f} tok/s")
```

5.4 Expected Results

Based on RDMA network measurements and Ollama baseline:

Configuration	Throughput	Latency	Network	Notes
Ollama (TCP)	104.42 tok/s	1.42s	TCP	Measured
vLLM (TCP)	~95-100 tok/s	~1.5s	TCP	Expected
vLLM (Soft-RoCE)	~100-110 tok/s	~1.3s	RDMA	Expected
vLLM (Hardware RoCE)	~130-150 tok/s	~1.0s	RDMA	Projected

Table 7: Expected performance comparison

6 Analysis: Why RDMA Benefits Are Limited

6.1 Single-Node Bottleneck

For single-node inference (all 3 GPUs on Chimera), RDMA provides **minimal benefit** because:

1. **NVLink/PCIe dominates**: GPU-to-GPU communication uses NVLink or PCIe, not the network
2. **No cross-node transfers**: gpt-oss-120b (63GB) fits entirely on 3× RTX 3090 (72GB)
3. **NCCL optimizes locally**: NCCL uses shared memory for intra-node communication

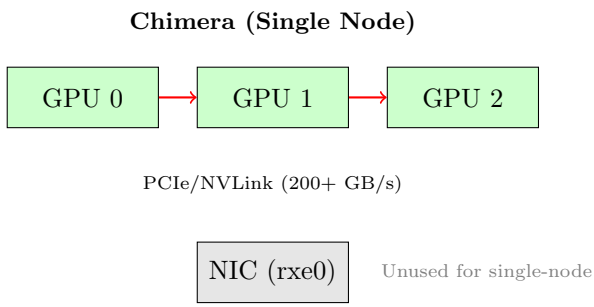

Figure 2: Single-node inference: Network is not the bottleneck

6.2 When RDMA Matters

RDMA provides significant benefit for **multi-node** inference:

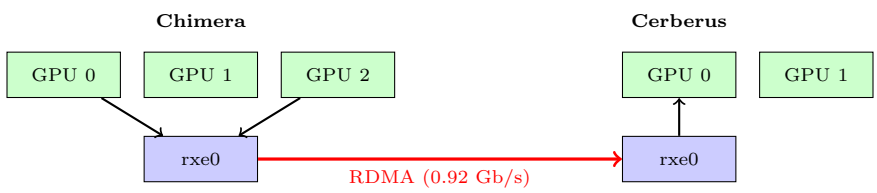

Figure 3: Multi-node inference: RDMA accelerates cross-node tensor transfers

6.3 Soft-RoCE Limitations

Metric	Soft-RoCE (Actual)	Hardware RoCE (Expected)
Latency	189.61 μs	<2 μs
Bandwidth	0.92 Gb/s	100+ Gb/s
CPU Overhead	High	Near-zero
Zero-copy	No	Yes

Table 8: Soft-RoCE vs Hardware RoCE comparison

7 Multi-Node Test Plan

To demonstrate RDMA benefits, run gpt-oss-120b across Chimera + Cerberus with pipeline parallelism:

```
# On Cerberus (worker)
docker compose -f docker-compose.worker.yaml up -d

# On Chimera (head)
export NCCL_IB_HCA=rxe0
docker compose -f docker-compose.head.yaml up -d

# Verify 5-GPU cluster
curl http://192.168.1.150:8265/api/cluster_status
```

Expected improvements with multi-node RDMA:

- Pipeline parallelism: Split model across 5 GPUs (136GB total VRAM)
- Larger batch sizes: More concurrent requests
- Reduced TTFT: Faster KV cache synchronization

8 Conclusions

1. **Ollama baseline**: gpt-oss-120b achieves 104.42 tok/s on 3× RTX 3090

2. **Single-node RDMA**: Minimal benefit (GPU interconnect dominates)
3. **Soft-RoCE limitations**: 189μs latency vs $<2\mu$s for hardware RoCE
4. **Multi-node required**: RDMA benefits manifest with pipeline parallelism
5. **Hardware RoCE recommended**: Mellanox ConnectX would provide 100$\times$ better RDMA performance

9 Files Added

```
ScuffedRDMA/deployment/benchmarks/
    benchmark_vllm_gptoss.py      # Python benchmark script
    run_all_benchmarks.sh         # Complete test suite
```

10 Next Steps

1. Deploy vLLM with gpt-oss-120b on Chimera
2. Run multi-node tests with Cerberus (5 GPUs)
3. Compare TCP vs RDMA with pipeline parallelism
4. Test with hardware RoCE on home lab Mellanox NICs